

COMP5318-assignment2-notebook(2)

May 13, 2025

1 COMP4318/5318 Assignment 2: Image Classification

1.0.1 Group number: 283 , SID1: 540282735 , SID2: 540114883

This template notebook includes code to load the dataset and a skeleton for the main sections that should be included in the notebook. Please stick to this struture for your submitted notebook.

Please focus on making your code clear, with appropriate variable names and whitespace. Include comments and markdown text to aid the readability of your code where relevant. See the specification and marking criteria in the associated specification to guide you when completing your implementation.

1.1 Setup and dependencies

Please use this section to list and set up all your required libraries/dependencies and your plotting environment.

```
[3]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV,
    StratifiedGroupKFold, RandomizedSearchCV
from sklearn.decomposition import PCA
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC

from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

import keras_tuner
import tensorflow as tf
from tensorflow import keras

from sklearn.model_selection import ParameterGrid
import pandas as pd
import time
```

1.2 1. Data loading, exploration, and preprocessing

Code to load the dataset is provided in the following cell. Please proceed with your data exploration and preprocessing in the remainder of this section.


```
[4]: # Load the dataset training and test sets as numpy arrays
# assuming Assignment2Data folder is present in the same directory
# as the notebook
X_train = np.load('Assignment2Data/X_train.npy')
y_train = np.load('Assignment2Data/y_train.npy')
X_test = np.load('Assignment2Data/X_test.npy')
y_test = np.load('Assignment2Data/y_test.npy')
```

```
[5]: print(X_train.shape)
```

(32000, 28, 28, 3)

```
[6]: def plot_examples(data, label, num_rows=4, num_cols=10):
    plt.figure(figsize=(num_cols * 1.2, num_rows * 1.5))
    for row in range(num_rows):
        for col in range(num_cols):
            index = num_cols * row + col
            plt.subplot(num_rows, num_cols, index + 1)
            plt.imshow(data[index])
            plt.title(f"Label: {label[index]}")
            plt.axis('off')
    plt.tight_layout()
    plt.show()

plot_examples(X_train, y_train)
```



```
[7]: print(np.min(X_train), np.max(X_train))

0 255

[8]: scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train.reshape(X_train.shape[0], -1)).
    ↪reshape(X_train.shape)
X_test_scaled = scaler.transform(X_test.reshape(X_test.shape[0], -1)).
    ↪reshape(X_test.shape)


[9]: X_train_flat = X_train.reshape(X_train.shape[0], -1)
X_test_flat = X_test.reshape(X_test.shape[0], -1)

[10]: print(X_train_flat.shape)

(32000, 2352)

[11]: pca = PCA()
pca.fit(X_train_flat)
cumsum = np.cumsum(pca.explained_variance_ratio_)

[12]: plt.figure(figsize=(10, 6))
plt.plot(cumsum, linewidth=3)
plt.axis([0, 1000, 0.8, 1])
plt.xlabel("Dimensions")
plt.ylabel("Explained Variance")
plt.grid(True)
plt.show()
```



```
[13]: print("Number of components to retain 95% variance: ", np.argmax(cumsum >= 0.95) + 1)
print("Number of components to retain 98% variance: ", np.argmax(cumsum >= 0.98) + 1)
print("Number of components to retain 99% variance: ", np.argmax(cumsum >= 0.99) + 1)
```

Number of components to retain 95% variance: 332

Number of components to retain 98% variance: 566

Number of components to retain 99% variance: 719

```
[14]: pca = PCA(n_components=719)
X_train_pca = pca.fit_transform(X_train_flat)
X_test_pca = pca.transform(X_test_flat)
```


```
[15]: print(X_train_pca.shape)
```

(32000, 719)

1.2.1 Examples of preprocessed data

Please print/display some examples of your preprocessed data here.

```
[16]: plot_examples(X_train_scaled, y_train)
```



```
[17]: print(X_train.shape)
print(X_test_pca.shape)
```

```
print(X_train[0])
```

```
(32000, 28, 28, 3)
(8000, 719)
[[[199 129 181]
 [191 127 174]
 [179 118 173]

 ...
 [209 136 184]
 [209 118 178]
 [205 122 180]]]

[[202 125 181]
 [194 127 182]
 [184 117 174]

 ...
 [194 137 182]
 [162 95 154]
 [156 89 150]]]

[[165 102 156]
 [182 108 167]
 [190 115 174]

 ...
 [202 138 188]
 [181 125 178]
 [163 102 159]]]

...
[[208 124 182]
 [206 128 184]
 [217 123 181]

 ...
 [217 128 185]
 [212 135 189]
 [207 137 190]]]

[[203 123 180]
 [194 127 182]
 [204 125 182]

 ...
 [207 135 187]
 [205 137 190]
 [204 132 187]]]

[[202 122 180]
 [189 126 180]]
```

```
[184 122 176]  
...  
[202 123 181]  
[191 124 181]  
[190 117 176]]]
```

```
[18]: print(np.unique(y_train))
```

```
[0 1 2 3 4 5 6 7 8]
```

1.3 2. Algorithm design and setup

1.3.1 Algorithm of choice from first six weeks of course

```
[19]: def build_svm(param):  
    model = SVC(**param)  
    return model
```

1.3.2 Fully connected neural network

```
[20]: def build_mlp(hp):  
    model = keras.Sequential([  
        keras.layers.Input(shape=(X_train_pca.shape[1],)),  
        keras.layers.Dense(units=hp.Int('units_1', min_value=128,  
        max_value=1024, step=64), activation='relu'), # Updated range  
        keras.layers.Dense(units=hp.Int('units_2', min_value=64, max_value=512,  
        step=32), activation='relu'), # Updated range  
        keras.layers.Dense(9, activation='softmax')  
    ])  
  
    model.compile(  
        optimizer=keras.optimizers.Adam(hp.Choice('learning_rate',  
        values=[5e-4, 1e-3, 5e-3])), # Updated values  
        loss='sparse_categorical_crossentropy',  
        metrics=['accuracy'])  
    )  
  
    return model
```

1.3.3 Convolutional neural network

```
[21]: def build_cnn(hp):  
    model = keras.Sequential([  
        keras.layers.Input((28, 28, 3)),  
        keras.layers.Conv2D(filters=hp.Int('conv_1_filters', min_value=32,  
        max_value=128, step=32),  
                           kernel_size=hp.Choice('conv_1_kernel', values=[3,  
                           5]),
```

```

        activation='relu', padding='same'),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Conv2D(filters=hp.Int('conv_2_filters', min_value=64,
    max_value=256, step=64),
                        kernel_size=hp.Choice('conv_2_kernel', values=[3,
    5]),
        activation='relu', padding='same'),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(units=hp.Int('dense_units', min_value=128,
    max_value=1024, step=128), activation='relu'),
    keras.layers.Dense(9, activation='softmax')
])

model.compile(
    optimizer=keras.optimizers.Adam(hp.Choice('learning_rate',
    values=[5e-4, 1e-3, 5e-3])),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

return model

```

1.4 3. Hyperparameter tuning

```
[22]: X_small, _, y_small, _ = train_test_split(
    X_train_pca, y_train,
    train_size=6000,
    stratify=y_train,
    random_state=42
)

X_train_val, X_val, y_train_val, y_val = train_test_split(
    X_small, y_small,
    test_size=0.2,
    stratify=y_small,
    random_state=42
)
```

1.4.1 Algorithm of choice from first six weeks of course

```
[23]: svm_param_grid = {
    'C': [1, 10],
    'kernel': ['rbf', 'linear', 'poly', 'sigmoid'],
    'gamma': ['scale']
}
```

```

svm_search = RandomizedSearchCV(
    estimator=SVC(),
    param_distributions=svm_param_grid,
    n_iter=5,
    cv=3,
    n_jobs=-1,
    verbose=2
)

svm_search.fit(X_train_val, y_train_val)

```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

```

[CV] END ...C=10, gamma=scale, kernel=sigmoid; total time= 6.0s
[CV] END ...C=10, gamma=scale, kernel=sigmoid; total time= 6.1s
[CV] END ...C=10, gamma=scale, kernel=sigmoid; total time= 6.0s
[CV] END ...C=10, gamma=scale, kernel=linear; total time= 7.6s
[CV] END ...C=10, gamma=scale, kernel=linear; total time= 7.6s
[CV] END ...C=10, gamma=scale, kernel=linear; total time= 7.8s
[CV] END ...C=10, gamma=scale, kernel=rbf; total time= 9.5s
[CV] END ...C=10, gamma=scale, kernel=rbf; total time= 9.6s
[CV] END ...C=1, gamma=scale, kernel=rbf; total time= 7.5s
[CV] END ...C=1, gamma=scale, kernel=rbf; total time= 7.6s
[CV] END ...C=10, gamma=scale, kernel=rbf; total time= 8.3s
[CV] END ...C=1, gamma=scale, kernel=linear; total time= 6.4s
[CV] END ...C=1, gamma=scale, kernel=linear; total time= 6.5s
[CV] END ...C=1, gamma=scale, kernel=rbf; total time= 7.0s
[CV] END ...C=1, gamma=scale, kernel=linear; total time= 5.4s

```

[23]: RandomizedSearchCV(cv=3, estimator=SVC(), n_iter=5, n_jobs=-1,
 param_distributions={'C': [1, 10], 'gamma': ['scale'],
 'kernel': ['rbf', 'linear', 'poly',
 'sigmoid']},
 verbose=2)

[24]: results_df = pd.DataFrame(svm_search.cv_results_)

results_df = results_df[[
 'params',
 'mean_test_score',
 'std_test_score',
 'rank_test_score'
]]

results_df = results_df.sort_values(by='mean_test_score', ascending=False)

print(results_df)

```

print("Best parameters: ", svm_search.best_params_)
results_df.to_csv("svm_search_results.csv", index=False)

            params  mean_test_score \
3      {'kernel': 'rbf', 'gamma': 'scale', 'C': 1}      0.591875
2      {'kernel': 'rbf', 'gamma': 'scale', 'C': 10}     0.578542
1      {'kernel': 'linear', 'gamma': 'scale', 'C': 10}    0.381250
4      {'kernel': 'linear', 'gamma': 'scale', 'C': 1}     0.381250
0      {'kernel': 'sigmoid', 'gamma': 'scale', 'C': 10}    0.282708

   std_test_score  rank_test_score
3        0.010206          1
2        0.010223          2
1        0.010607          3
4        0.010607          3
0        0.035858          5

Best parameters:  {'kernel': 'rbf', 'gamma': 'scale', 'C': 1}

```

1.4.2 Fully connected neural network

```
[25]: mlp_tuner = keras_tuner.RandomSearch(
    hypermodel=build_mlp,
    objective="val_accuracy",
    max_trials=10,  # Increased trials
    executions_per_trial=3,  # Increased executions
    overwrite=True,
    directory="result",
    project_name="mlp"
)

mlp_tuner.search(X_train_val, y_train_val, epochs=30, validation_data=(X_val, y_val))
```

Trial 10 Complete [00h 00m 33s]
val_accuracy: 0.42249999443689984

Best val_accuracy So Far: 0.5099999904632568
Total elapsed time: 00h 04m 16s

```
[26]: print("Best MLP parameters:", mlp_tuner.get_best_hyperparameters()[0].values)

trials = list(mlp_tuner.oracle.trials.values())

results = [
    **trial.hyperparameters.values,
    "val_accuracy": trial.metrics.get_last_value("val_accuracy"),
    "score": trial.score
} for trial in trials]
```

```
results_df = pd.DataFrame(results)
results_df = results_df.sort_values(by="val_accuracy", ascending=False)
```

```
print(results_df)
```

```
Best MLP parameters: {'units_1': 448, 'units_2': 480, 'learning_rate': 0.0005}
   units_1  units_2  learning_rate  val_accuracy      score
7      448      480       0.0005    0.510000  0.510000
2      320      352       0.0010    0.506111  0.506111
4      256      320       0.0010    0.496944  0.496944
8      960      416       0.0010    0.495833  0.495833
3      128      256       0.0010    0.485000  0.485000
0      128      192       0.0005    0.461111  0.461111
1      192       96       0.0005    0.454722  0.454722
5      704      320       0.0050    0.435833  0.435833
9      640      512       0.0050    0.422500  0.422500
6     1024      288       0.0050    0.367222  0.367222
```

1.4.3 Convolutional neural network

```
[27]: cnn_tuner = keras_tuner.RandomSearch(
    hypermodel=build_cnn,
    objective='val_accuracy',
    max_trials=10,
    executions_per_trial=3,
    overwrite=True,
    directory="result",
    project_name="cnn"
)
```

```
[28]: X_tune, _, y_tune, _ = train_test_split(X_train_scaled, y_train, □
    ↪train_size=6000, stratify=y_train, random_state=42)

cnn_tuner.search(X_tune, y_tune, epochs=30, validation_split=0.2)
```

```
Trial 10 Complete [00h 22m 15s]
val_accuracy: 0.7322222391764323
```

```
Best val_accuracy So Far: 0.7463888923327128
Total elapsed time: 02h 23m 30s
```

```
[29]: print("Best CNN parameters:", cnn_tuner.get_best_hyperparameters()[0].values)

cnn_trials = list(cnn_tuner.oracle.trials.values())
```

```

cnn_results = [
    **trial.hyperparameters.values,
    "val_accuracy": trial.metrics.get_last_value("val_accuracy"),
    "score": trial.score
} for trial in cnn_trials]

cnn_results_df = pd.DataFrame(cnn_results)
cnn_results_df = cnn_results_df.sort_values(by="val_accuracy", ascending=False)

print(cnn_results_df)

```

Best CNN parameters: {'conv_1_filters': 32, 'conv_1_kernel': 3, 'conv_2_filters': 128, 'conv_2_kernel': 3, 'dense_units': 768, 'learning_rate': 0.0005}

	conv_1_filters	conv_1_kernel	conv_2_filters	conv_2_kernel	dense_units	\
6	32	3	128	3	768	
8	128	3	64	3	640	
0	32	3	192	3	256	
9	128	3	128	3	256	
2	32	3	64	5	768	
3	32	5	64	3	640	
1	96	5	64	5	256	
7	64	5	64	3	128	
4	64	5	256	3	512	
5	32	5	256	3	512	
	learning_rate	val_accuracy	score			
6	0.0005	0.746389	0.746389			
8	0.0005	0.743889	0.743889			
0	0.0010	0.735000	0.735000			
9	0.0010	0.732222	0.732222			
2	0.0010	0.716667	0.716667			
3	0.0005	0.713889	0.713889			
1	0.0005	0.702500	0.702500			
7	0.0005	0.700000	0.700000			
4	0.0010	0.690833	0.690833			
5	0.0050	0.419444	0.419444			

1.5 4. Final models

In this section, please ensure to include cells to train each model with its best hyperparameter combination independently of the hyperparameter tuning cells, i.e. don't rely on the hyperparameter tuning cells having been run.

1.5.1 Algorithm of choice from first six weeks of course

```
[30]: # the best parameters
best_svm = SVC(kernel='rbf', C=1, gamma='scale')

# train
start_time = time.time()
best_svm.fit(X_train_pca, y_train)
svm_train_time = time.time() - start_time
print(f"SVM training time: {svm_train_time:.2f} seconds")

svm_accuracy = best_svm.score(X_test_pca, y_test)
print(f"Best SVM Test Accuracy: {svm_accuracy:.3f}")
```

SVM training time: 237.36 seconds

Best SVM Test Accuracy: 0.659


```
[31]: y_pred_svm = best_svm.predict(X_test_pca)

print("SVM Classification Report:")
print(classification_report(y_test, y_pred_svm))

plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_svm), annot=True, fmt='d',
            cmap='Blues')
plt.title("SVM Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

SVM Classification Report:

	precision	recall	f1-score	support
0	0.83	0.95	0.89	873
1	0.96	0.88	0.91	858
2	0.47	0.30	0.36	877
3	0.91	0.85	0.88	914
4	0.58	0.58	0.58	737
5	0.64	0.68	0.66	1072
6	0.44	0.26	0.32	682
7	0.48	0.52	0.50	813
8	0.55	0.77	0.64	1174
accuracy			0.66	8000
macro avg	0.65	0.64	0.64	8000
weighted avg	0.65	0.66	0.65	8000



1.5.2 Fully connected neural network

```
[32]: # best parameter: units_1=576, units_2=448, learning_rate=0.001
mlp = keras.Sequential([
    keras.layers.Input(shape=(X_train_pca.shape[1],)),
    keras.layers.Dense(576, activation='relu'),
    keras.layers.Dense(448, activation='relu'),
    keras.layers.Dense(9, activation='softmax')
])

mlp.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001),
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

start_time = time.time()
mlp.fit(X_train_pca, y_train, epochs=100, batch_size=64)
mlp_train_time = time.time() - start_time
print(f"MLP training time: {mlp_train_time:.2f} seconds")

_, mlp_accuracy = mlp.evaluate(X_test_pca, y_test)
print(f"Best MLP Test Accuracy: {mlp_accuracy:.3f}")
```

Epoch 1/100

```
500/500          2s 3ms/step -  
accuracy: 0.4111 - loss: 14.5633  
Epoch 2/100  
500/500          1s 3ms/step -  
accuracy: 0.6173 - loss: 1.0120  
Epoch 3/100  
500/500          1s 3ms/step -  
accuracy: 0.7014 - loss: 0.8001  
Epoch 4/100  
500/500          1s 3ms/step -  
accuracy: 0.7591 - loss: 0.6459  
Epoch 5/100  
500/500          1s 3ms/step -  
accuracy: 0.7959 - loss: 0.5663  
Epoch 6/100  
500/500          1s 2ms/step -  
accuracy: 0.8141 - loss: 0.5157  
Epoch 7/100  
500/500          1s 3ms/step -  
accuracy: 0.8226 - loss: 0.5021  
Epoch 8/100  
500/500          1s 2ms/step -  
accuracy: 0.8356 - loss: 0.4823  
Epoch 9/100  
500/500          1s 3ms/step -  
accuracy: 0.8468 - loss: 0.4596  
Epoch 10/100  
500/500          1s 2ms/step -  
accuracy: 0.8588 - loss: 0.4093  
Epoch 11/100  
500/500          1s 3ms/step -  
accuracy: 0.8611 - loss: 0.4113  
Epoch 12/100  
500/500          1s 3ms/step -  
accuracy: 0.8613 - loss: 0.4173  
Epoch 13/100  
500/500          1s 3ms/step -  
accuracy: 0.8864 - loss: 0.3382  
Epoch 14/100  
500/500          1s 3ms/step -  
accuracy: 0.8891 - loss: 0.3373  
Epoch 15/100  
500/500          1s 3ms/step -  
accuracy: 0.8933 - loss: 0.3227  
Epoch 16/100  
500/500          1s 3ms/step -  
accuracy: 0.8954 - loss: 0.3263  
Epoch 17/100
```

```
500/500          1s 3ms/step -  
accuracy: 0.9015 - loss: 0.3147  
Epoch 18/100  
500/500          1s 3ms/step -  
accuracy: 0.9116 - loss: 0.2858  
Epoch 19/100  
500/500          1s 3ms/step -  
accuracy: 0.9135 - loss: 0.2718  
Epoch 20/100  
500/500          1s 3ms/step -  
accuracy: 0.9110 - loss: 0.2848  
Epoch 21/100  
500/500          1s 3ms/step -  
accuracy: 0.9159 - loss: 0.2744  
Epoch 22/100  
500/500          1s 3ms/step -  
accuracy: 0.9208 - loss: 0.2566  
Epoch 23/100  
500/500          1s 3ms/step -  
accuracy: 0.9219 - loss: 0.2689  
Epoch 24/100  
500/500          1s 3ms/step -  
accuracy: 0.9168 - loss: 0.2858  
Epoch 25/100  
500/500          1s 3ms/step -  
accuracy: 0.9270 - loss: 0.2428  
Epoch 26/100  
500/500          1s 3ms/step -  
accuracy: 0.9268 - loss: 0.2457  
Epoch 27/100  
500/500          1s 3ms/step -  
accuracy: 0.9387 - loss: 0.2053  
Epoch 28/100  
500/500          1s 3ms/step -  
accuracy: 0.9384 - loss: 0.1982  
Epoch 29/100  
500/500          1s 3ms/step -  
accuracy: 0.9299 - loss: 0.2461  
Epoch 30/100  
500/500          1s 3ms/step -  
accuracy: 0.9317 - loss: 0.2414  
Epoch 31/100  
500/500          1s 3ms/step -  
accuracy: 0.9394 - loss: 0.2083  
Epoch 32/100  
500/500          1s 3ms/step -  
accuracy: 0.9404 - loss: 0.2055  
Epoch 33/100
```

```
500/500          1s 3ms/step -  
accuracy: 0.9242 - loss: 0.2942  
Epoch 34/100  
500/500          1s 3ms/step -  
accuracy: 0.9401 - loss: 0.2308  
Epoch 35/100  
500/500          1s 3ms/step -  
accuracy: 0.9411 - loss: 0.2382  
Epoch 36/100  
500/500          1s 3ms/step -  
accuracy: 0.9473 - loss: 0.2017  
Epoch 37/100  
500/500          1s 3ms/step -  
accuracy: 0.9508 - loss: 0.1737  
Epoch 38/100  
500/500          1s 3ms/step -  
accuracy: 0.9475 - loss: 0.1970  
Epoch 39/100  
500/500          1s 3ms/step -  
accuracy: 0.9487 - loss: 0.1926  
Epoch 40/100  
500/500          1s 3ms/step -  
accuracy: 0.9446 - loss: 0.2187  
Epoch 41/100  
500/500          1s 3ms/step -  
accuracy: 0.9491 - loss: 0.1929  
Epoch 42/100  
500/500          1s 3ms/step -  
accuracy: 0.9508 - loss: 0.1779  
Epoch 43/100  
500/500          1s 3ms/step -  
accuracy: 0.9536 - loss: 0.1742  
Epoch 44/100  
500/500          1s 3ms/step -  
accuracy: 0.9507 - loss: 0.1926  
Epoch 45/100  
500/500          1s 3ms/step -  
accuracy: 0.9494 - loss: 0.1942  
Epoch 46/100  
500/500          1s 3ms/step -  
accuracy: 0.9549 - loss: 0.1929  
Epoch 47/100  
500/500          1s 3ms/step -  
accuracy: 0.9559 - loss: 0.1776  
Epoch 48/100  
500/500          1s 3ms/step -  
accuracy: 0.9517 - loss: 0.2248  
Epoch 49/100
```

```
500/500          1s 3ms/step -  
accuracy: 0.9595 - loss: 0.1717  
Epoch 50/100  
500/500          1s 3ms/step -  
accuracy: 0.9562 - loss: 0.1611  
Epoch 51/100  
500/500          1s 3ms/step -  
accuracy: 0.9554 - loss: 0.1948  
Epoch 52/100  
500/500          1s 3ms/step -  
accuracy: 0.9593 - loss: 0.1670  
Epoch 53/100  
500/500          1s 3ms/step -  
accuracy: 0.9533 - loss: 0.2023  
Epoch 54/100  
500/500          1s 3ms/step -  
accuracy: 0.9622 - loss: 0.1607  
Epoch 55/100  
500/500          1s 3ms/step -  
accuracy: 0.9573 - loss: 0.1881  
Epoch 56/100  
500/500          1s 3ms/step -  
accuracy: 0.9577 - loss: 0.1912  
Epoch 57/100  
500/500          1s 3ms/step -  
accuracy: 0.9567 - loss: 0.2051  
Epoch 58/100  
500/500          1s 3ms/step -  
accuracy: 0.9592 - loss: 0.1729  
Epoch 59/100  
500/500          1s 3ms/step -  
accuracy: 0.9559 - loss: 0.1931  
Epoch 60/100  
500/500          1s 3ms/step -  
accuracy: 0.9578 - loss: 0.1841  
Epoch 61/100  
500/500          1s 3ms/step -  
accuracy: 0.9611 - loss: 0.1749  
Epoch 62/100  
500/500          1s 3ms/step -  
accuracy: 0.9685 - loss: 0.1504  
Epoch 63/100  
500/500          1s 3ms/step -  
accuracy: 0.9613 - loss: 0.1762  
Epoch 64/100  
500/500          1s 3ms/step -  
accuracy: 0.9641 - loss: 0.1537  
Epoch 65/100
```

```
500/500          1s 3ms/step -  
accuracy: 0.9650 - loss: 0.1504  
Epoch 66/100  
500/500          1s 3ms/step -  
accuracy: 0.9552 - loss: 0.2292  
Epoch 67/100  
500/500          1s 3ms/step -  
accuracy: 0.9597 - loss: 0.1914  
Epoch 68/100  
500/500          1s 3ms/step -  
accuracy: 0.9648 - loss: 0.1534  
Epoch 69/100  
500/500          1s 3ms/step -  
accuracy: 0.9600 - loss: 0.2259  
Epoch 70/100  
500/500          1s 3ms/step -  
accuracy: 0.9571 - loss: 0.2217  
Epoch 71/100  
500/500          1s 3ms/step -  
accuracy: 0.9662 - loss: 0.1491  
Epoch 72/100  
500/500          1s 3ms/step -  
accuracy: 0.9663 - loss: 0.1560  
Epoch 73/100  
500/500          1s 3ms/step -  
accuracy: 0.9565 - loss: 0.2122  
Epoch 74/100  
500/500          1s 3ms/step -  
accuracy: 0.9645 - loss: 0.1786  
Epoch 75/100  
500/500          1s 3ms/step -  
accuracy: 0.9681 - loss: 0.1536  
Epoch 76/100  
500/500          1s 3ms/step -  
accuracy: 0.9636 - loss: 0.1748  
Epoch 77/100  
500/500          1s 3ms/step -  
accuracy: 0.9623 - loss: 0.1892  
Epoch 78/100  
500/500          1s 3ms/step -  
accuracy: 0.9674 - loss: 0.1497  
Epoch 79/100  
500/500          1s 3ms/step -  
accuracy: 0.9679 - loss: 0.1532  
Epoch 80/100  
500/500          1s 3ms/step -  
accuracy: 0.9648 - loss: 0.1724  
Epoch 81/100
```

```
500/500          1s 3ms/step -  
accuracy: 0.9588 - loss: 0.2350  
Epoch 82/100  
500/500          2s 3ms/step -  
accuracy: 0.9649 - loss: 0.1778  
Epoch 83/100  
500/500          1s 3ms/step -  
accuracy: 0.9668 - loss: 0.1702  
Epoch 84/100  
500/500          1s 3ms/step -  
accuracy: 0.9713 - loss: 0.1293  
Epoch 85/100  
500/500          1s 3ms/step -  
accuracy: 0.9633 - loss: 0.1961  
Epoch 86/100  
500/500          1s 3ms/step -  
accuracy: 0.9642 - loss: 0.1752  
Epoch 87/100  
500/500          1s 3ms/step -  
accuracy: 0.9683 - loss: 0.1596  
Epoch 88/100  
500/500          1s 3ms/step -  
accuracy: 0.9646 - loss: 0.2019  
Epoch 89/100  
500/500          1s 3ms/step -  
accuracy: 0.9690 - loss: 0.2252  
Epoch 90/100  
500/500          1s 3ms/step -  
accuracy: 0.9689 - loss: 0.1564  
Epoch 91/100  
500/500          1s 3ms/step -  
accuracy: 0.9679 - loss: 0.1601  
Epoch 92/100  
500/500          1s 3ms/step -  
accuracy: 0.9698 - loss: 0.1635  
Epoch 93/100  
500/500          1s 3ms/step -  
accuracy: 0.9675 - loss: 0.1750  
Epoch 94/100  
500/500          1s 3ms/step -  
accuracy: 0.9711 - loss: 0.1490  
Epoch 95/100  
500/500          1s 3ms/step -  
accuracy: 0.9679 - loss: 0.1710  
Epoch 96/100  
500/500          1s 3ms/step -  
accuracy: 0.9696 - loss: 0.1661  
Epoch 97/100
```

```

500/500           1s 3ms/step -
accuracy: 0.9655 - loss: 0.2215
Epoch 98/100
500/500           1s 3ms/step -
accuracy: 0.9653 - loss: 0.1888
Epoch 99/100
500/500           1s 3ms/step -
accuracy: 0.9707 - loss: 0.1536
Epoch 100/100
500/500           1s 3ms/step -
accuracy: 0.9730 - loss: 0.1553
MLP training time: 137.65 seconds
250/250          0s 625us/step -
accuracy: 0.5974 - loss: 9.7343
Best MLP Test Accuracy: 0.597

```

```

[33]: y_pred_mlp = mlp.predict(X_test_pca)
y_pred_labels_mlp = y_pred_mlp.argmax(axis=1)

print("MLP Classification Report:")
print(classification_report(y_test, y_pred_labels_mlp))

plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_labels_mlp), annot=True, fmt='d', cmap='Greens')
plt.title("MLP Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()


```

```

250/250          0s 509us/step
MLP Classification Report:
      precision    recall   f1-score   support
          0       0.81     0.90     0.86     873
          1       0.92     0.90     0.91     858
          2       0.42     0.39     0.40     877
          3       0.82     0.75     0.78     914
          4       0.49     0.48     0.48     737
          5       0.59     0.60     0.59    1072
          6       0.31     0.42     0.36     682
          7       0.41     0.38     0.39     813
          8       0.56     0.51     0.53    1174

   accuracy          0.60     8000
macro avg       0.59     0.59     8000
weighted avg    0.60     0.60     8000

```



1.5.3 Convolutional neural network

```
[34]: # best parameter: {'conv_1_filters': 96, 'conv_1_kernel': 3, 'conv_2_filters': 64, 'conv_2_kernel': 3, 'dense_units': 640, 'learning_rate': 0.0005}
cnn = keras.Sequential([
    keras.layers.Input(shape=(28, 28, 3)),
    keras.layers.Conv2D(96, kernel_size=3, activation='relu', padding='same'),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Conv2D(64, kernel_size=3, activation='relu', padding='same'),
    keras.layers.MaxPooling2D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(640, activation='relu'),
    keras.layers.Dense(9, activation='softmax')
])

cnn.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005),
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

start_time = time.time()
```

```

cnn.fit(X_train_scaled, y_train, epochs=100, batch_size=64)
cnn_train_time = time.time() - start_time
print(f"CNN training time: {cnn_train_time:.2f} seconds")

_, cnn_accuracy = cnn.evaluate(X_test_scaled, y_test)
print(f"Best CNN Test Accuracy: {cnn_accuracy:.3f}")

```

Epoch 1/100
500/500 18s 35ms/step -
accuracy: 0.3873 - loss: 1.6017
Epoch 2/100
500/500 17s 35ms/step -
accuracy: 0.6598 - loss: 0.9202
Epoch 3/100
500/500 18s 35ms/step -
accuracy: 0.7303 - loss: 0.7413
Epoch 4/100
500/500 18s 35ms/step -
accuracy: 0.7662 - loss: 0.6451
Epoch 5/100
500/500 18s 35ms/step -
accuracy: 0.7886 - loss: 0.5802
Epoch 6/100
500/500 18s 35ms/step -
accuracy: 0.8110 - loss: 0.5272
Epoch 7/100
500/500 18s 36ms/step -
accuracy: 0.8186 - loss: 0.4969
Epoch 8/100
500/500 18s 36ms/step -
accuracy: 0.8322 - loss: 0.4633
Epoch 9/100
500/500 18s 36ms/step -
accuracy: 0.8469 - loss: 0.4202
Epoch 10/100
500/500 18s 37ms/step -
accuracy: 0.8487 - loss: 0.4152
Epoch 11/100
500/500 18s 36ms/step -
accuracy: 0.8632 - loss: 0.3808
Epoch 12/100
500/500 18s 36ms/step -
accuracy: 0.8731 - loss: 0.3472
Epoch 13/100
500/500 18s 35ms/step -
accuracy: 0.8902 - loss: 0.3115
Epoch 14/100

```
500/500          18s 35ms/step -  
accuracy: 0.8959 - loss: 0.2868  
Epoch 15/100  
500/500          18s 35ms/step -  
accuracy: 0.9002 - loss: 0.2762  
Epoch 16/100  
500/500          18s 36ms/step -  
accuracy: 0.9105 - loss: 0.2482  
Epoch 17/100  
500/500          18s 36ms/step -  
accuracy: 0.9197 - loss: 0.2235  
Epoch 18/100  
500/500          18s 35ms/step -  
accuracy: 0.9251 - loss: 0.2102  
Epoch 19/100  
500/500          18s 35ms/step -  
accuracy: 0.9280 - loss: 0.1980  
Epoch 20/100  
500/500          18s 35ms/step -  
accuracy: 0.9422 - loss: 0.1628  
Epoch 21/100  
500/500          18s 36ms/step -  
accuracy: 0.9470 - loss: 0.1496  
Epoch 22/100  
500/500          18s 35ms/step -  
accuracy: 0.9498 - loss: 0.1429  
Epoch 23/100  
500/500          18s 35ms/step -  
accuracy: 0.9571 - loss: 0.1224  
Epoch 24/100  
500/500          18s 35ms/step -  
accuracy: 0.9562 - loss: 0.1277  
Epoch 25/100  
500/500          18s 36ms/step -  
accuracy: 0.9606 - loss: 0.1124  
Epoch 26/100  
500/500          18s 36ms/step -  
accuracy: 0.9695 - loss: 0.0927  
Epoch 27/100  
500/500          18s 36ms/step -  
accuracy: 0.9782 - loss: 0.0716  
Epoch 28/100  
500/500          18s 36ms/step -  
accuracy: 0.9768 - loss: 0.0705  
Epoch 29/100  
500/500          18s 36ms/step -  
accuracy: 0.9833 - loss: 0.0564  
Epoch 30/100
```

```
500/500          18s 36ms/step -  
accuracy: 0.9842 - loss: 0.0512  
Epoch 31/100  
500/500          18s 36ms/step -  
accuracy: 0.9804 - loss: 0.0602  
Epoch 32/100  
500/500          18s 36ms/step -  
accuracy: 0.9866 - loss: 0.0454  
Epoch 33/100  
500/500          18s 36ms/step -  
accuracy: 0.9859 - loss: 0.0447  
Epoch 34/100  
500/500          18s 36ms/step -  
accuracy: 0.9891 - loss: 0.0420  
Epoch 35/100  
500/500          18s 36ms/step -  
accuracy: 0.9759 - loss: 0.0699  
Epoch 36/100  
500/500          18s 36ms/step -  
accuracy: 0.9866 - loss: 0.0436  
Epoch 37/100  
500/500          18s 36ms/step -  
accuracy: 0.9947 - loss: 0.0218  
Epoch 38/100  
500/500          18s 36ms/step -  
accuracy: 0.9881 - loss: 0.0397  
Epoch 39/100  
500/500          18s 36ms/step -  
accuracy: 0.9780 - loss: 0.0689  
Epoch 40/100  
500/500          18s 36ms/step -  
accuracy: 0.9941 - loss: 0.0219  
Epoch 41/100  
500/500          18s 36ms/step -  
accuracy: 0.9883 - loss: 0.0361  
Epoch 42/100  
500/500          18s 36ms/step -  
accuracy: 0.9909 - loss: 0.0289  
Epoch 43/100  
500/500          18s 36ms/step -  
accuracy: 0.9855 - loss: 0.0458  
Epoch 44/100  
500/500          18s 36ms/step -  
accuracy: 0.9899 - loss: 0.0306  
Epoch 45/100  
500/500          18s 36ms/step -  
accuracy: 0.9842 - loss: 0.0520  
Epoch 46/100
```

```
500/500          18s 36ms/step -  
accuracy: 0.9911 - loss: 0.0299  
Epoch 47/100  
500/500          18s 36ms/step -  
accuracy: 0.9885 - loss: 0.0363  
Epoch 48/100  
500/500          18s 36ms/step -  
accuracy: 0.9925 - loss: 0.0261  
Epoch 49/100  
500/500          18s 36ms/step -  
accuracy: 0.9843 - loss: 0.0485  
Epoch 50/100  
500/500          18s 36ms/step -  
accuracy: 0.9974 - loss: 0.0107  
Epoch 51/100  
500/500          18s 36ms/step -  
accuracy: 0.9827 - loss: 0.0502  
Epoch 52/100  
500/500          18s 36ms/step -  
accuracy: 0.9984 - loss: 0.0075  
Epoch 53/100  
500/500          18s 37ms/step -  
accuracy: 0.9931 - loss: 0.0225  
Epoch 54/100  
500/500          18s 36ms/step -  
accuracy: 0.9926 - loss: 0.0260  
Epoch 55/100  
500/500          18s 36ms/step -  
accuracy: 0.9980 - loss: 0.0082  
Epoch 56/100  
500/500          18s 36ms/step -  
accuracy: 0.9862 - loss: 0.0449  
Epoch 57/100  
500/500          18s 36ms/step -  
accuracy: 0.9985 - loss: 0.0076  
Epoch 58/100  
500/500          18s 36ms/step -  
accuracy: 0.9855 - loss: 0.0445  
Epoch 59/100  
500/500          18s 36ms/step -  
accuracy: 0.9945 - loss: 0.0160  
Epoch 60/100  
500/500          18s 36ms/step -  
accuracy: 0.9893 - loss: 0.0322  
Epoch 61/100  
500/500          18s 36ms/step -  
accuracy: 0.9992 - loss: 0.0038  
Epoch 62/100
```

```
500/500          18s 36ms/step -  
accuracy: 0.9965 - loss: 0.0115  
Epoch 63/100  
500/500          18s 36ms/step -  
accuracy: 0.9906 - loss: 0.0281  
Epoch 64/100  
500/500          18s 36ms/step -  
accuracy: 0.9961 - loss: 0.0138  
Epoch 65/100  
500/500          18s 36ms/step -  
accuracy: 0.9959 - loss: 0.0132  
Epoch 66/100  
500/500          18s 37ms/step -  
accuracy: 0.9984 - loss: 0.0067  
Epoch 67/100  
500/500          18s 36ms/step -  
accuracy: 0.9813 - loss: 0.0596  
Epoch 68/100  
500/500          18s 36ms/step -  
accuracy: 0.9998 - loss: 0.0012  
Epoch 69/100  
500/500          18s 37ms/step -  
accuracy: 0.9998 - loss: 0.0012  
Epoch 70/100  
500/500          18s 37ms/step -  
accuracy: 0.9947 - loss: 0.0173  
Epoch 71/100  
500/500          18s 37ms/step -  
accuracy: 0.9985 - loss: 0.0080  
Epoch 72/100  
500/500          18s 37ms/step -  
accuracy: 0.9961 - loss: 0.0150  
Epoch 73/100  
500/500          18s 37ms/step -  
accuracy: 0.9874 - loss: 0.0368  
Epoch 74/100  
500/500          18s 37ms/step -  
accuracy: 0.9994 - loss: 0.0033  
Epoch 75/100  
500/500          18s 37ms/step -  
accuracy: 0.9985 - loss: 0.0067  
Epoch 76/100  
500/500          18s 36ms/step -  
accuracy: 0.9905 - loss: 0.0310  
Epoch 77/100  
500/500          18s 36ms/step -  
accuracy: 1.0000 - loss: 9.7236e-04  
Epoch 78/100
```

```
500/500          18s 36ms/step -  
accuracy: 1.0000 - loss: 4.2210e-04  
Epoch 79/100  
500/500          18s 36ms/step -  
accuracy: 0.9790 - loss: 0.0649  
Epoch 80/100  
500/500          18s 36ms/step -  
accuracy: 0.9938 - loss: 0.0226  
Epoch 81/100  
500/500          18s 36ms/step -  
accuracy: 0.9995 - loss: 0.0021  
Epoch 82/100  
500/500          18s 36ms/step -  
accuracy: 0.9912 - loss: 0.0306  
Epoch 83/100  
500/500          18s 36ms/step -  
accuracy: 0.9999 - loss: 0.0012  
Epoch 84/100  
500/500          18s 36ms/step -  
accuracy: 0.9991 - loss: 0.0034  
Epoch 85/100  
500/500          18s 36ms/step -  
accuracy: 0.9876 - loss: 0.0406  
Epoch 86/100  
500/500          18s 36ms/step -  
accuracy: 0.9997 - loss: 0.0021  
Epoch 87/100  
500/500          18s 37ms/step -  
accuracy: 0.9998 - loss: 0.0013  
Epoch 88/100  
500/500          18s 37ms/step -  
accuracy: 0.9779 - loss: 0.0729  
Epoch 89/100  
500/500          18s 37ms/step -  
accuracy: 0.9995 - loss: 0.0027  
Epoch 90/100  
500/500          19s 38ms/step -  
accuracy: 0.9899 - loss: 0.0307  
Epoch 91/100  
500/500          19s 37ms/step -  
accuracy: 0.9996 - loss: 0.0024  
Epoch 92/100  
500/500          18s 37ms/step -  
accuracy: 0.9972 - loss: 0.0094  
Epoch 93/100  
500/500          18s 37ms/step -  
accuracy: 0.9948 - loss: 0.0181  
Epoch 94/100
```

```

500/500          18s 36ms/step -
accuracy: 0.9998 - loss: 0.0018
Epoch 95/100
500/500          18s 37ms/step -
accuracy: 0.9994 - loss: 0.0029
Epoch 96/100
500/500          18s 37ms/step -
accuracy: 0.9875 - loss: 0.0367
Epoch 97/100
500/500          18s 37ms/step -
accuracy: 0.9981 - loss: 0.0071
Epoch 98/100
500/500          18s 37ms/step -
accuracy: 0.9989 - loss: 0.0052
Epoch 99/100
500/500          19s 37ms/step -
accuracy: 0.9875 - loss: 0.0371
Epoch 100/100
500/500          18s 37ms/step -
accuracy: 0.9999 - loss: 9.8921e-04
CNN training time: 1810.31 seconds
250/250          1s 5ms/step -
accuracy: 0.8536 - loss: 0.9121
Best CNN Test Accuracy: 0.857

```

```
[35]: y_pred_cnn = cnn.predict(X_test_scaled)
y_pred_labels_cnn = y_pred_cnn.argmax(axis=1)

print("CNN Classification Report:")
print(classification_report(y_test, y_pred_labels_cnn))

plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_labels_cnn), annot=True, fmt='d', cmap='Oranges')
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

```

250/250          1s 5ms/step
CNN Classification Report:
      precision    recall  f1-score   support

       0       0.94      0.99      0.96      873
       1       0.98      0.97      0.98      858
       2       0.81      0.71      0.76      877

```

3	0.95	0.96	0.95	914
4	0.83	0.87	0.85	737
5	0.83	0.84	0.83	1072
6	0.80	0.79	0.80	682
7	0.67	0.68	0.67	813
8	0.86	0.88	0.87	1174
accuracy			0.86	8000
macro avg	0.85	0.85	0.85	8000
weighted avg	0.86	0.86	0.86	8000

